

# High Volume Data Feed (HVDF)

- open-source framework created by the team at MongoDB
- makes it easy to efficiently validate, store, index, query and purge time series data in MongoDB via a simple REST API

# HVDF Incoming Data Has Three Major Components

- Feed: Maps to a database
- Channel: Maps to a collection
- Sample: Maps to a document

# HVDF Enables

- Use MongoDB in an append-only fashion
- Store documents sequentially
- Apply time-based partitioning

# Time Partitioning

- A new collection at a specified time interval
- Sequential writes
- Fast deletes by dropping a collection
- Reduced Index size
- Optimized collections for reads
- Automatic sharding: HVDF automatically creates a shard key for each time interval as the collection is created

# HVDF Architecture

